Social Media Sentiments

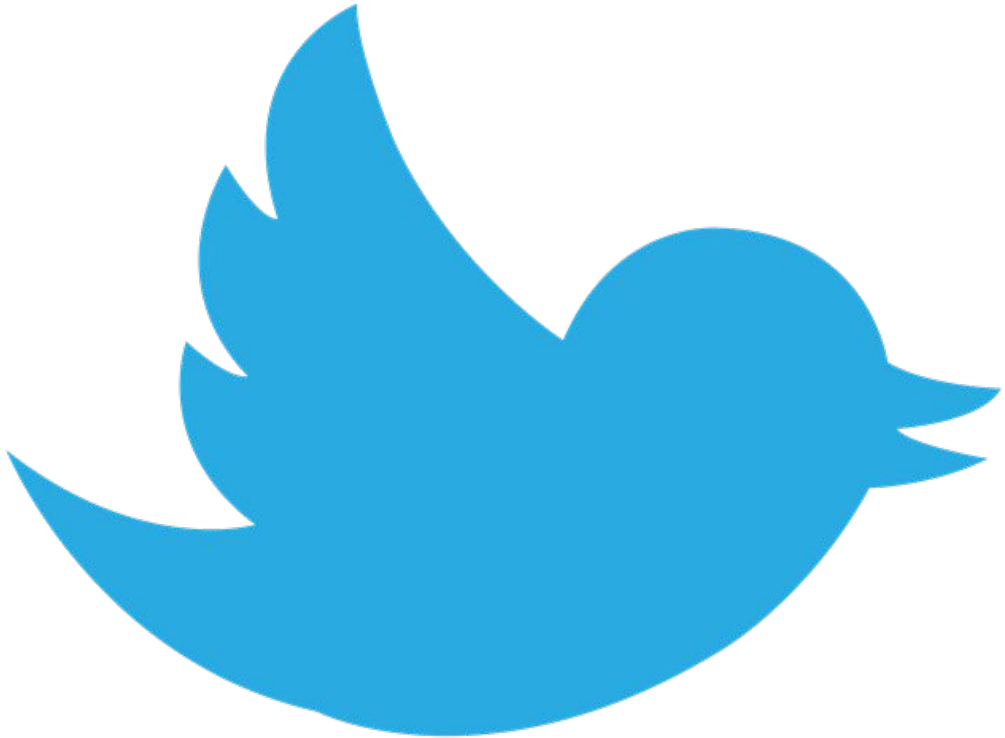

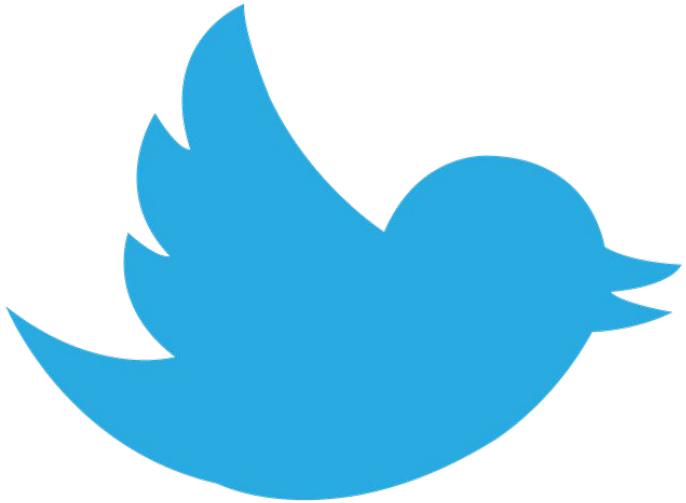

Twitter often is blamed for furmenting ill feelings on its user base. I wanted to know: at any given moment, what is the likelihood for a tweet to be either racist or misogynistic rather than positive and uplifting.

In a collection of nearly 32,000 tweets, I’ve picked up keywords that reflect either positive or negative sentiments, and then with machine learning built an algorithm to predict the likelihood a tweet projecting negativity

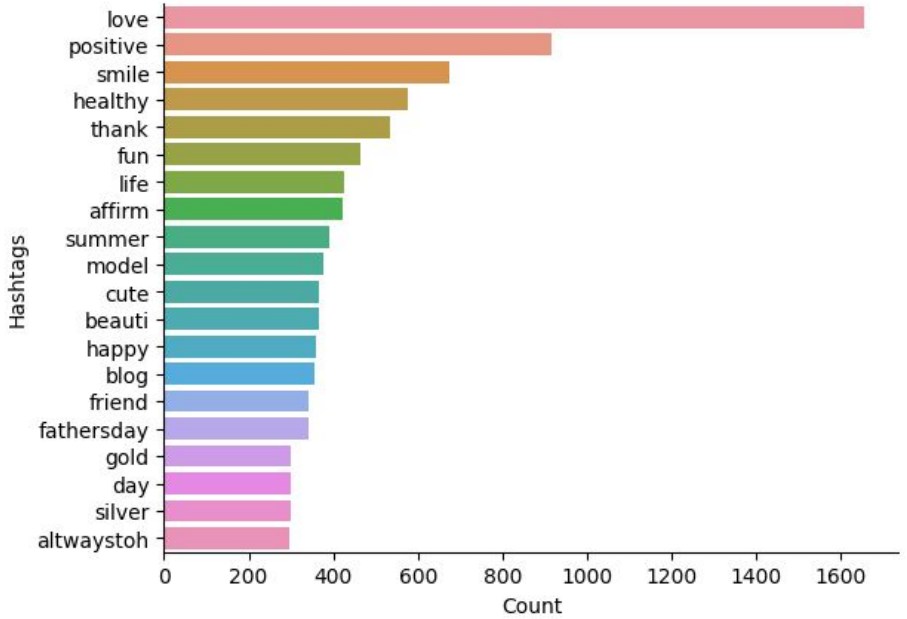

It’s not unexpected that *love* would top the list as the most commonly used keyword to convey positivity. However, it is remarkable that the word *positive* takes a distant second place.

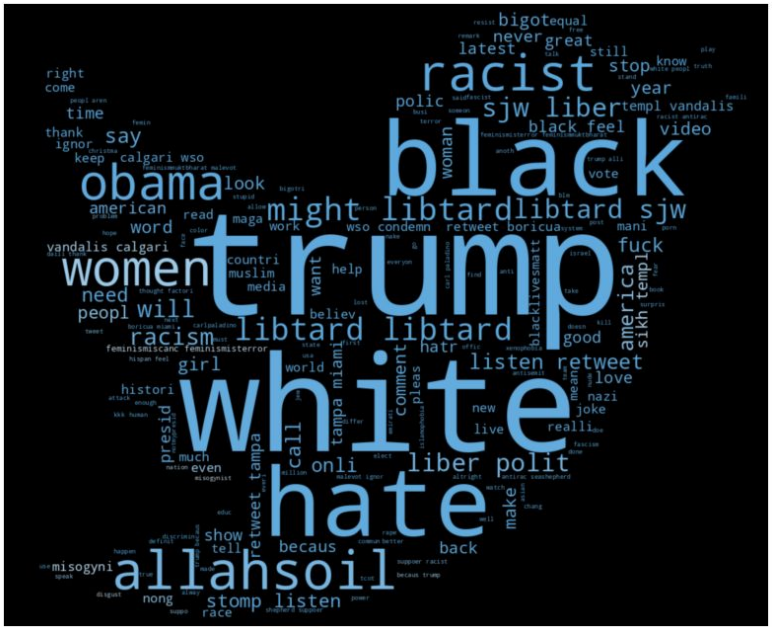

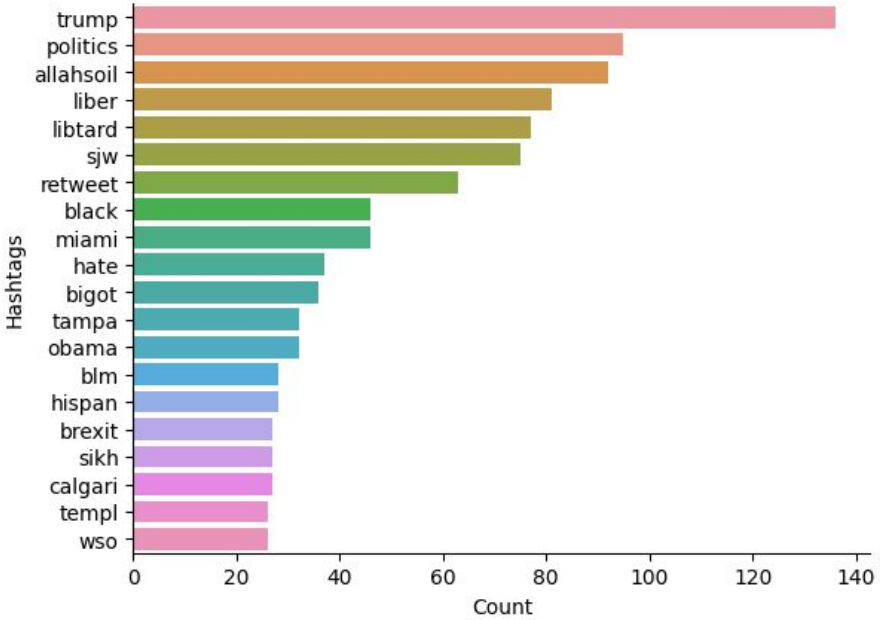

Political keywords abound! With the current heightened political tension, it's unsurprising to find representation of racial tensions between *black* and *white*, along with the most popular political figure of our era.

F1 Score

Training algorithms LogisticRegression, XGBoost, and DecisionTree on feature text data Bag-of-Words and TF-IDF. The closest number to 1 indicates the highest level of accuracy.

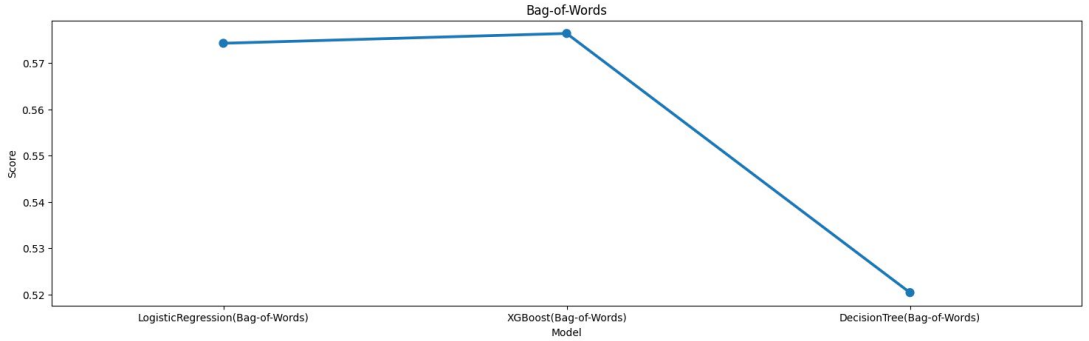

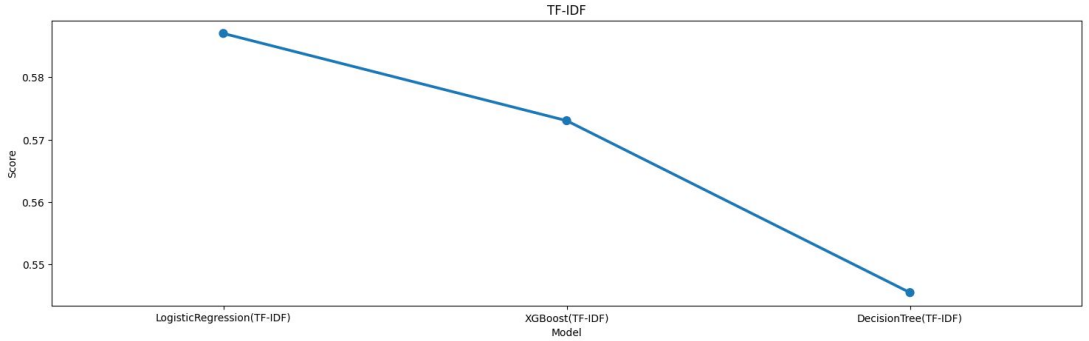

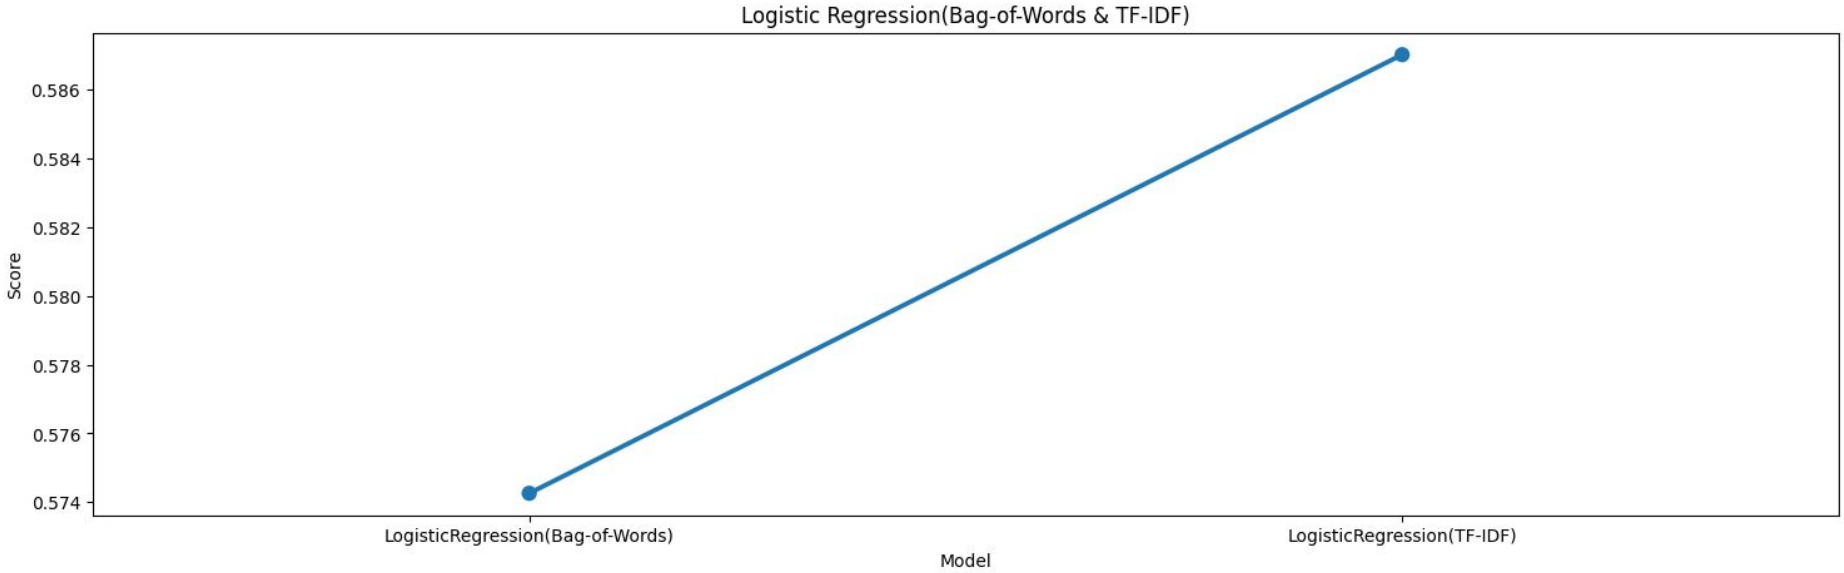

The LogisticRegression model trained on TF-IDF data is the most accurate and will be employed as a predictor for the test data

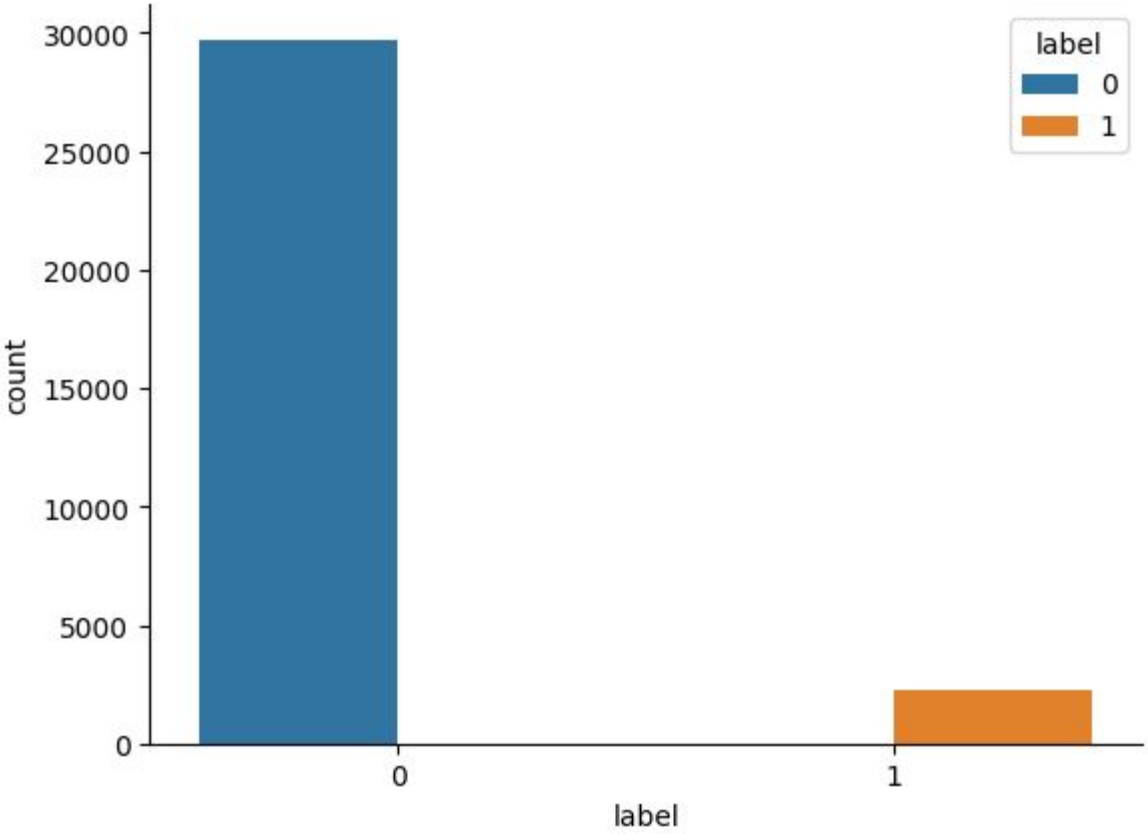

It is **1225.602%** more likely for a tweet to be negative with sexist or racist content than positive with non-sexist or non-racist content